AtliQ PromiseKeeper: AI Opportunity Map

Root cause from Step 1: AtliQ does not have a lead problem or a discipline problem. It has a memory problem: revenue leaks because commitments die in inboxes. $423k+ of open pipeline has an unanswered client or a deadline within 3 weeks (as of 10 Jul 2026).
Read top to bottom: 1 Workflow, 2 Score, 3 Prioritise, 4 Scope, 5 Alternatives rejected, 6 Feature to AI capability and guardrails.

1. Founder-seller workflow: where promises leak, and where AI can help

MVP | Later (v2) | Not doing

Six stages of a founder-led deal. Red: today's leak, with dataset evidence. Chips: AI opportunities, scored in frames 2 and 3.

1 Capture the lead

TODAY'S LEAK

Leads arrive by email intro, phone or business card. Entering them is a separate step that gets skipped.

EVIDENCE IN THE DATASET

Harbor and Finch, Nimbus Lending and PixelWorks are not in the CRM. About 10 TechSummit cards sit in a drawer. Ridgeline Foods called and left no contact details.

WHERE AI CAN HELP

- **O4** Lead and referral capture
- **O11** Business-card scanning

2 Qualify and discover

TODAY'S LEAK

Call notes are scribbles or forwarded to Jay. Budget and decision-maker facts never reach the CRM.

EVIDENCE IN THE DATASET

AgroLink revealed a 35-40k budget, still blank in CRM. Dhaval on MiraHealth: I'll forget half of it by next week. EduSpark named a new decision-maker, unlogged.

WHERE AI CAN HELP

- **O9** Account history Q and A before calls
- **O14** Call recording and transcription

3 Propose and negotiate

TODAY'S LEAK

After a proposal, the client's questions and dated asks land in one inbox and go unseen.

EVIDENCE IN THE DATASET

Meridian's two IT questions unanswered since 16 Jun. ShopSwift needs an SOW by 15 Jul. VoltEnergy asked for a 3-year TCO on 1 Jul, no reply.

WHERE AI CAN HELP

- **O2** Client deadline and unanswered-request radar

4 Follow up and keep promises

TODAY'S LEAK

I'll send it by Friday is typed on a phone between meetings, then forgotten. Nobody owns the next step.

EVIDENCE IN THE DATASET

13 promises made by AtliQ were not kept (ByteBridge, CloudNine, Zenlytics, Northwind, Acme). The NovaPharma questionnaire has no owner. 7 Jul meeting: no owners assigned.

WHERE AI CAN HELP

- **O1** Promise tracker
- **O8** Suggested owner for every item
- **O3** Daily ranked brief per seller
- **O10** Draft follow-up emails
- **O13** Auto-send follow-ups

5 Close and record

TODAY'S LEAK

Wins, approvals and status changes are not recorded. The CRM shows stale stages and duplicates.

EVIDENCE IN THE DATASET

Trustline is approved but still Proposal Sent. DataNest renewal agreed, CRM unchanged. Meridian and GlobalMart each appear twice. 0 of 40 records have a next follow-up date.

WHERE AI CAN HELP

- **O7** CRM fix suggestions for approval
- **O12** Win-probability lead scoring

6 Grow and revisit

TODAY'S LEAK

Expansion asks and check back later are mentioned in passing, then lost.

EVIDENCE IN THE DATASET

6 existing clients signalled new work (Acme, GlobalMart, Northwind, Sunrise, PayTrack, DataNest). Quantia's Q2 check-back was missed. Kaveri and ClearPath came back but are still Lost.

WHERE AI CAN HELP

- **O5** Cross-sell and expansion signals
- **O6** Revisit and lost-but-active alerts

2. Score every opportunity: impact vs feasibility

Scores 1 to 10. Feasibility is the average of data, build and adoption, the same method as the course Feasibility-Impact sheet.

Opportunity scores (sorted by score)

ID	Opportunity	Data	Build	Adoption	Feasibility (avg)	Impact	Score (F x I)	Decision
O2	Client deadline and unanswered-request radar	9	8	9	8.7	10	87.0	MVP
O1	Promise tracker (I'll send it by Friday)	9	7	9	8.3	10	83.0	MVP
O3	Daily ranked brief per seller	8	9	9	8.7	9	78.3	MVP
O4	Lead and referral capture	8	7	8	7.7	9	69.3	MVP
O5	Cross-sell and expansion signals	8	7	7	7.3	8	58.4	MVP
O8	Suggested owner for every item	7	8	6	7.0	8	56.0	MVP
O6	Revisit and lost-but-active alerts	8	8	8	8.0	7	56.0	MVP
O7	CRM fix suggestions (status, owner, duplicates)	7	6	8	7.0	7	49.0	MVP
O9	Account history Q and A before calls	7	7	6	6.7	6	40.2	Later (v2)
O10	Draft follow-up emails (human sends)	8	8	6	7.3	5	36.5	Later (v2)
O13	Auto-send follow-ups and auto-write CRM	8	7	2	5.7	6	34.2	Not doing
O11	Business-card scanning	3	6	7	5.3	4	21.2	Not now
O14	Call recording and transcription	1	4	3	2.7	6	16.2	Not now
O12	Win-probability lead scoring	2	5	3	3.3	4	13.2	Not doing

How we scored
Data: is the signal actually present in the emails, notes and CRM? (Call recordings and card images are not, so O14 and O11 score low.)
Build: can we build it reliably for the capstone with an LLM plus simple rules?
Adoption: will founder-sellers use it with zero data entry? Anything that acts on its own, or that needs typing, scores low.
Impact: revenue at stake in the dataset, plus how many of the six failure patterns from Step 1 it addresses.
Decision rule: MVP = Quick Wins quadrant and directly fixes the root cause. Later = useful, but does not fix remembering. Not now / Not doing = no data, too little data, or breaks trust.

3. Prioritise: impact vs feasibility

Each dot is an opportunity from frame 2. Midpoint 5 on both axes.

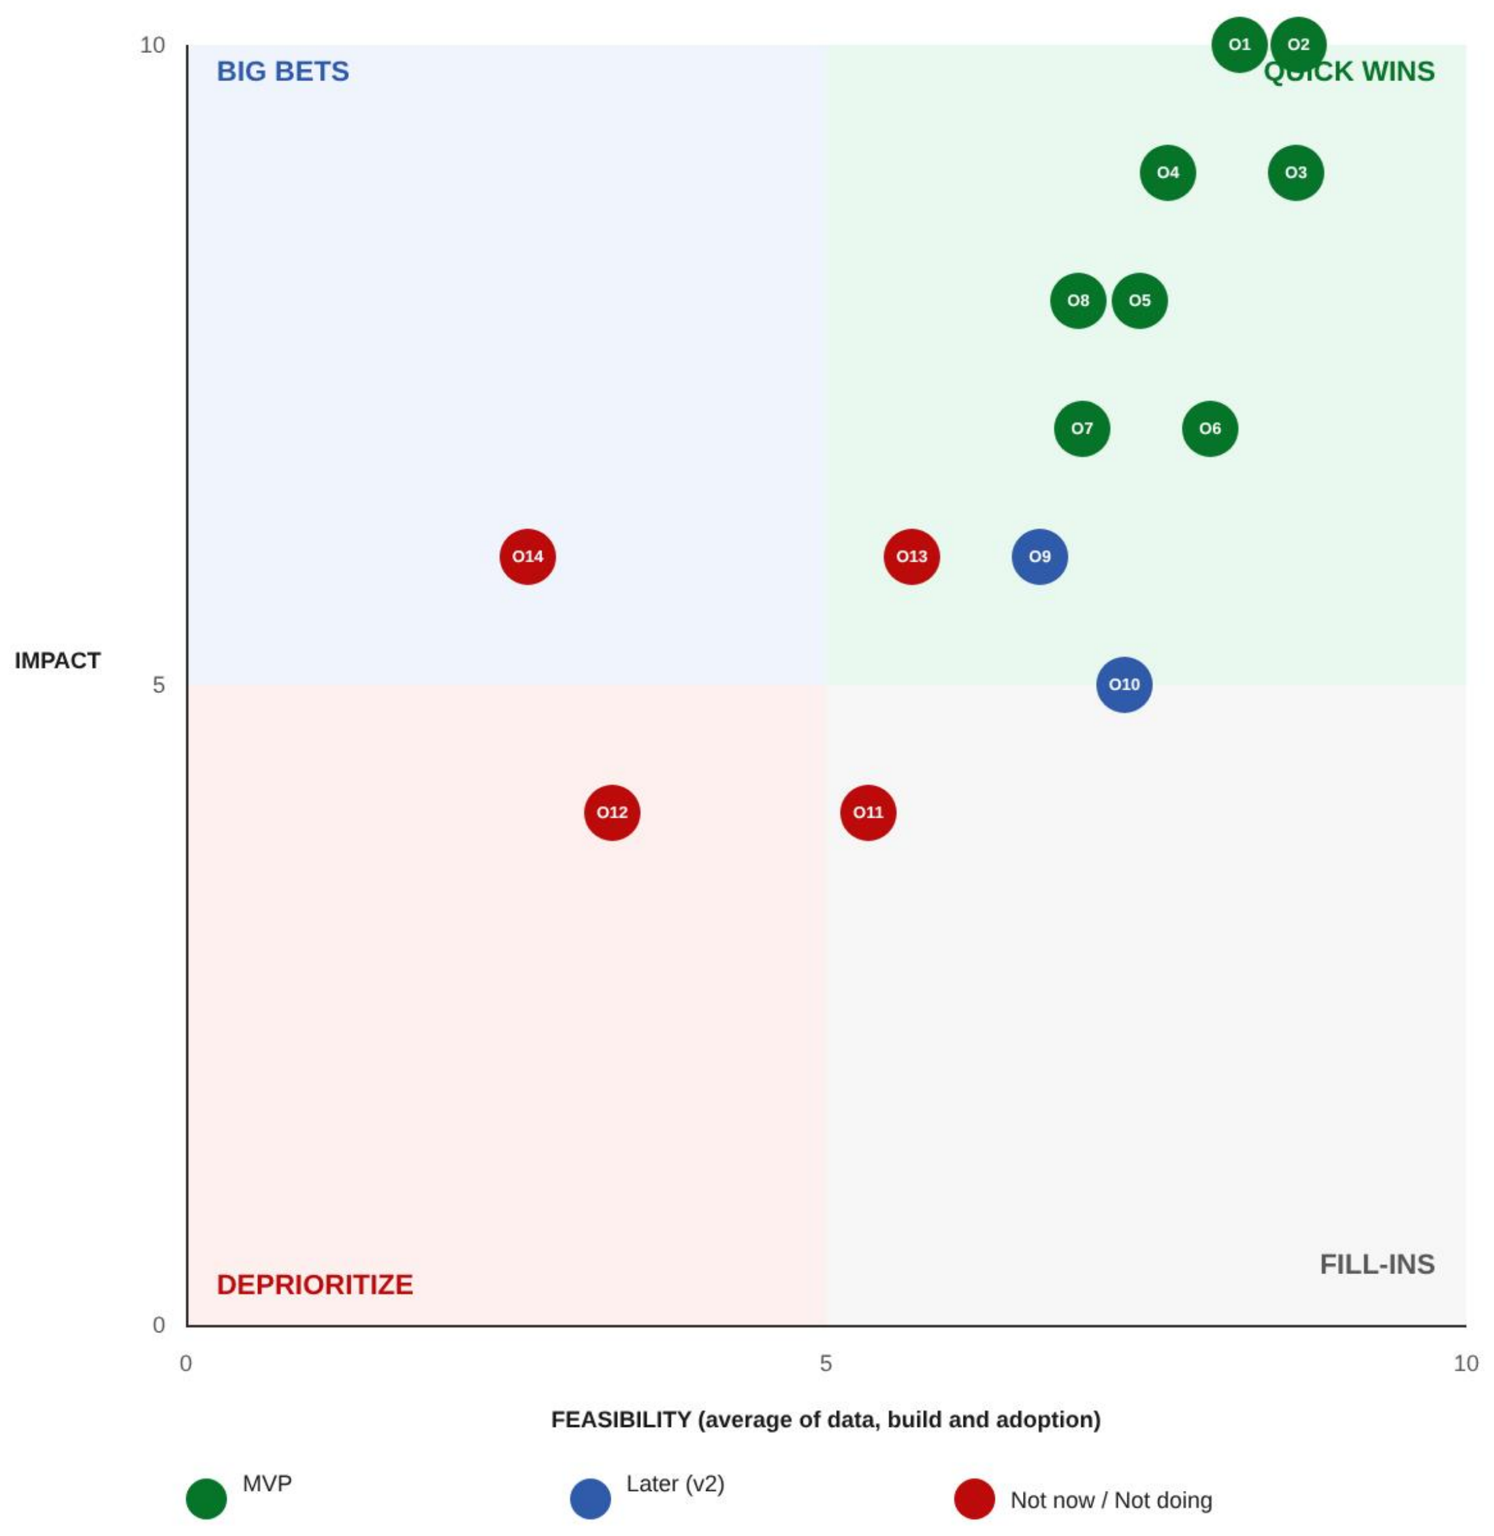

4. Scope: what PromiseKeeper should and should NOT do

The right column is as much a product decision as the left. Each line becomes a guardrail in frame 6 and in the PRD.

PROMISEKEEPER SHOULD	PROMISEKEEPER SHOULD NOT
Read emails, meeting notes and the CRM; extract commitments, client deadlines and buying signals	Send any email or message to a client
Attach the exact source quote and file to every item, or drop the item	Write to the CRM without a human approving the change
Suggest an owner and a due date for every item	Merge duplicates automatically (GlobalMart may be two real deals)
Rank each seller's daily brief with simple, visible rules: deadline, deal value, relationship	Predict win probability or score leads (5 wins in 40 records is too little data)
Propose CRM fixes, new leads and duplicate merges for one-tap human approval	Invent a fact, value or date that is not in a source
Say it is unsure and ask when evidence is ambiguous (e.g. three different Kulkarnis)	Follow instructions written inside an email (prompt injection)
	Rank or judge the sellers' performance: it is a memory aid, not surveillance
	Touch delivery or project work (out of scope per the dataset README)

5. Alternatives we considered and rejected

Why PromiseKeeper, and not the obvious first idea.

Auto-fill the CRM from email	Fixes the symptom. Records would be fuller, but nobody would see what is owed, to whom, or by when.
Lead scoring / win probability	AtliQ does not lack leads. And 40 records with 5 wins is far too little data to train or trust a score.
Email-drafting assistant	Makes sending faster, not remembering to send. Kept as a v2 add-on (O10) once promises are tracked.
Business-card scanner	Fixes one leak (about 10 cards) out of six failure patterns. No card images exist to test it.
More CRM discipline and reminders	Pledged in April and again in July; it failed both times. Any fix that needs typing will fail a third time.
Plain rules, no AI	Rules cannot read 'push in July after their board meet'. We keep rules for dates and ranking, and use AI only to read text.
Just use ChatGPT	No memory across three sources, no citations, no approval flow, no guardrails against sending or inventing facts.
PromiseKeeper (chosen)	Fixes the root cause: captures commitments with zero data entry, uses AI only where text must be understood, and a human approves every change.

6. Each MVP feature mapped to an AI capability and its guardrails

The 8 MVP opportunities group into 6 features. AI reads text; rules handle dates and ranking; a human approves every change.

PromiseKeeper MVP: feature to AI capability to guardrails

Feature	Opportunities	AI capability	Input (context)	Output	Guardrails	Human in the loop
F1 Promise tracker	O1, O8	LLM extraction: commitment detection, who / what / to whom, relative-date normalisation	AtliQ's sent emails and meeting notes	Commitment card: owner, client, what was promised, due date, source quote	No item without a verbatim quote; relative dates resolved against a set 'today' and the original phrase shown; low confidence goes to an Unsure list	Owner marks done, snoozes or dismisses
F2 Deadline and request radar	O2	LLM classification: open client request vs answered; date extraction	Inbound client threads (editorial notes stripped before the AI sees them)	Client deadline or unanswered request, with days remaining	Thread-level check that the last message is from the client; never infers a deadline that is not stated	Owner confirms or dismisses
F3 Daily brief	O3	Deterministic ranking rules plus LLM summarisation	Items from F1, F2, F4, F5 and CRM deal values	At most 10 ranked items per seller, each with a one-line reason and a source link	Ranking rules are visible (no black-box score); cap on items to avoid alert fatigue; summaries must cite sources	Seller can reorder, snooze or reassign
F4 Lead and referral capture	O4	Entity extraction plus entity resolution against the CRM	Emails, notes and CRM contacts	Proposed new CRM record with contact, source and referrer	Match on email and company; ambiguous match means ask, never guess; never creates a record by itself	One-tap approve, edit or reject
F5 Signal radar	O5, O6	LLM classification: cross-sell, revisit, lost-but-active	Client emails, QBR notes, lost-deal threads	Suggested opportunity or revisit date, with the client's own words	Must quote the client; no speculative cross-sell ideas; revisit dates only from what the client said	Account owner accepts or rejects
F6 CRM fix suggestions	O7	Comparison of CRM fields vs extracted facts; duplicate detection	CRM export plus outputs of F1 to F5	Suggested status, owner, follow-up date, value, or merge	Never auto-merges; shows both records side by side; every applied change is logged	Approve each change individually
Across all features	All	Guardrails layer	Every input	Audit log of every suggestion and decision	Email text treated as data, never as instructions; only needed fields sent to the model; precision and recall measured against our answer key before launch	Kill switch if any client contact or unapproved CRM write occurs